

KMeans Clustering with Facebook Thailand Data

Divij Pherwani

Contents

Initialization and Data Exploration	1
Data Engineering	4
Correlation Matrix	7
KMeans Clustering	9
Finding Optimal Number of Clusters	14
Analysis of $K = 2, 3$ and 4 . KMeans & PAM	18
Summary	33

Clustering is an unsupervised learning algorithm which groups data points into clusters. The set of data points in a group should have similar properties and/or features. There are several types of clustering methods. In this project I will be using K-means clustering which is also one of the most popular clustering algorithm. We identify K clusters of n observations that are grouped to their nearest centroid and each cluster will have a centroid containing the data points or vectors closest to it making centroid the center of each cluster.

More detailed theoretical introduction on K-Means clustering can be found here, <https://medium.com/0xcode/the-k-means-clustering-algorithm-intuition-demonstrated-in-r-aa62584a3649>.

In this project, I will be using a dataset from UCI Machine Learning Repository. The dataset contains Facebook pages of 10 Thai fashion and cosmetics retail sellers. Posts of a different nature (video, photos, statuses, and links). Engagement metrics consist of comments, shares, and reactions. The dataset has attributes such as `status_id`, `status_type`, `status_published`, `num_reactions`, `num_comments`, `num_shares`, `num_likes`, `num_loves`, `num_wows`, `num_hahas`, `num_sads` and `num_angrys`. The link to data set is <https://archive.ics.uci.edu/ml/datasets/Facebook+Live+Sellers+in+Thailand>.

Initialization and Data Exploration

```
# Loading libraries and setting directory
set.seed(100)
library(dplyr)
library(ggplot2)
library(cluster)
library(factoextra)
library(purrr)
library(tidyverse)
library(data.table)
library(corrplot)
library(flexclust)
```

```
library(fpc)
library(clustertend)
library(ClusterR)
library(stats)

setwd(dirname(rstudioapi::getSourceEditorContext()$path))
```

```
#Loading the dataset
data <- read.csv("Live_20210128.csv", stringsAsFactors = FALSE, header = TRUE)
head(data)
```

```
##      status_id status_type status_published num_reactions num_comments num_shares
## 1           1      video   4/22/2018 6:00           529           512           262
## 2           2      photo   4/21/2018 22:45           150             0             0
## 3           3      video   4/21/2018 6:17           227           236           57
## 4           4      photo   4/21/2018 2:29           111             0             0
## 5           5      photo   4/18/2018 3:22           213             0             0
## 6           6      photo   4/18/2018 2:14           217             6             0
##      num_likes num_loves num_wows num_hahas num_sads num_angrys Column1 Column2
## 1          432         92         3         1         1         0        NA        NA
## 2          150          0         0         0         0         0        NA        NA
## 3          204         21         1         1         0         0        NA        NA
## 4          111          0         0         0         0         0        NA        NA
## 5          204          9         0         0         0         0        NA        NA
## 6          211          5         1         0         0         0        NA        NA
##      Column3 Column4
## 1          NA        NA
## 2          NA        NA
## 3          NA        NA
## 4          NA        NA
## 5          NA        NA
## 6          NA        NA
```

```
#checking duplicates
data[duplicated(data),]
```

```
## [1] status_id      status_type      status_published num_reactions
## [5] num_comments    num_shares      num_likes        num_loves
## [9] num_wows        num_hahas       num_sads         num_angrys
## [13] Column1         Column2         Column3          Column4
## <0 rows> (or 0-length row.names)
```

```
# Summaary of dataset
summary(data)
```

```
##      status_id      status_type      status_published num_reactions
## Min.   : 1      Length:7050      Length:7050      Min.   : 0.0
## 1st Qu.:1763    Class :character  Class :character  1st Qu.: 17.0
## Median :3526    Mode  :character  Mode  :character  Median : 59.5
## Mean   :3526                                     Mean   : 230.1
## 3rd Qu.:5288                                     3rd Qu.: 219.0
## Max.   :7050                                     Max.   :4710.0
```

```
##      num_comments      num_shares      num_likes      num_loves
## Min.      :    0.0  Min.      :    0.00  Min.      :    0.0  Min.      :    0.00
## 1st Qu.:    0.0  1st Qu.:    0.00  1st Qu.:   17.0  1st Qu.:    0.00
## Median :    4.0  Median :    0.00  Median :   58.0  Median :    0.00
## Mean   :   224.4  Mean   :   40.02  Mean   :  215.0  Mean   :   12.73
## 3rd Qu.:   23.0  3rd Qu.:    4.00  3rd Qu.:  184.8  3rd Qu.:    3.00
## Max.    :20990.0  Max.    :3424.00  Max.    :4710.0  Max.    :657.00
##      num_wows      num_hahas      num_sads      num_angrys
## Min.      : 0.000  Min.      : 0.0000  Min.      : 0.0000  Min.      : 0.0000
## 1st Qu.: 0.000  1st Qu.: 0.0000  1st Qu.: 0.0000  1st Qu.: 0.0000
## Median : 0.000  Median : 0.0000  Median : 0.0000  Median : 0.0000
## Mean   :  1.289  Mean   :  0.6965  Mean   :  0.2437  Mean   :  0.1132
## 3rd Qu.: 0.000  3rd Qu.: 0.0000  3rd Qu.: 0.0000  3rd Qu.: 0.0000
## Max.    :278.000  Max.    :157.0000  Max.    :51.0000  Max.    :31.0000
## Column1      Column2      Column3      Column4
## Mode:logical  Mode:logical  Mode:logical  Mode:logical
## NA's:7050     NA's:7050     NA's:7050     NA's:7050
##
##
##
##
```

#Checking for NAs

```
head(data[,colSums(is.na(data)) > 0])
```

```
##      Column1 Column2 Column3 Column4
## 1         NA      NA      NA      NA
## 2         NA      NA      NA      NA
## 3         NA      NA      NA      NA
## 4         NA      NA      NA      NA
## 5         NA      NA      NA      NA
## 6         NA      NA      NA      NA
```

#remove last 4 columns

```
data <- data[,1:12]
```

```
data$status_published <- format(data$status_published, format = "%m/%d/%Y %H:%M")
```

#remove time and keeping date

```
data$status_published <- as.POSIXct(data$status_published, format = "%m/%d/%Y")
```

```
head(data)
```

```
##      status_id status_type status_published num_reactions num_comments num_shares
## 1           1      video    2018-04-22             529          512          262
## 2           2      photo    2018-04-21             150           0           0
## 3           3      video    2018-04-21             227          236           57
## 4           4      photo    2018-04-21             111           0           0
## 5           5      photo    2018-04-18             213           0           0
## 6           6      photo    2018-04-18             217           6           0
##      num_likes num_loves num_wows num_hahas num_sads num_angrys
## 1          432          92         3         1         1         0
## 2          150           0         0         0         0         0
## 3          204          21         1         1         0         0
## 4          111           0         0         0         0         0
## 5          204           9         0         0         0         0
## 6          211           5         1         0         0         0
```

```
#Checking for NAs after processing
data[rowSums(is.na(data)) > 0,]
```

```
## [1] status_id      status_type      status_published num_reactions
## [5] num_comments     num_shares      num_likes       num_loves
## [9] num_wows         num_hahas       num_sads        num_angrys
## <0 rows> (or 0-length row.names)
```

```
#Number of distinct rows
nrow(distinct(data))
```

```
## [1] 7050
```

```
# Unique status_type, we have 4 types of status - video, photo, link and status.
unique(data$status_type)
```

```
## [1] "video" "photo" "link"  "status"
```

```
# Checking the frequency of status update per day, it doesn't give any significant insight
x <- as.data.frame(plyr::count(data, "status_published"))
x <- x %>% arrange(desc(freq))
head(x)
```

```
##   status_published freq
## 1      2018-06-07   51
## 2      2018-06-09   43
## 3      2018-06-11   42
## 4      2018-05-25   38
## 5      2018-06-08   35
## 6      2018-05-23   33
```

```
# Analyszing internal structure of dataset
str(data)
```

```
## 'data.frame': 7050 obs. of 12 variables:
## $ status_id : int 1 2 3 4 5 6 7 8 9 10 ...
## $ status_type : chr "video" "photo" "video" "photo" ...
## $ status_published: POSIXct, format: "2018-04-22" "2018-04-21" ...
## $ num_reactions : int 529 150 227 111 213 217 503 295 203 170 ...
## $ num_comments : int 512 0 236 0 0 6 614 453 1 9 ...
## $ num_shares : int 262 0 57 0 0 0 72 53 0 1 ...
## $ num_likes : int 432 150 204 111 204 211 418 260 198 167 ...
## $ num_loves : int 92 0 21 0 9 5 70 32 5 3 ...
## $ num_wows : int 3 0 1 0 0 1 10 1 0 0 ...
## $ num_hahas : int 1 0 1 0 0 0 2 1 0 0 ...
## $ num_sads : int 1 0 0 0 0 0 0 0 0 0 ...
## $ num_angrys : int 0 0 0 0 0 0 3 1 0 0 ...
```

Data Engineering

```
# We don't need status_id, status_published date for building clusters so we drop them

data <- data[,c(2,4:12)]
head(data)
```

```
##   status_type num_reactions num_comments num_shares num_likes num_loves
## 1      video          529           512          262          432          92
## 2      photo          150            0            0          150            0
## 3      video          227           236            57          204           21
## 4      photo          111            0            0          111            0
## 5      photo          213            0            0          204            9
## 6      photo          217             6            0          211            5
##   num_wows num_hahas num_sads num_angrys
## 1         3         1         1         0
## 2         0         0         0         0
## 3         1         1         0         0
## 4         0         0         0         0
## 5         0         0         0         0
## 6         1         0         0         0
```

```
df <- data[,2:ncol(data)]
head(df)
```

```
##   num_reactions num_comments num_shares num_likes num_loves num_wows num_hahas
## 1           529           512          262          432          92          3          1
## 2           150            0            0          150            0          0          0
## 3           227           236            57          204           21          1          1
## 4           111            0            0          111            0          0          0
## 5           213            0            0          204            9          0          0
## 6           217             6            0          211            5          1          0
##   num_sads num_angrys
## 1         1         0
## 2         0         0
## 3         0         0
## 4         0         0
## 5         0         0
## 6         0         0
```

```
# we check the assumption that num_reactions = num_likes + num_loves + num_wows + num_hahas + num_sads

count <- rep(NA,nrow(df))
l <- rep(0,nrow(df))

for (i in 1:nrow(df))
{
  if(df[i,"num_reactions"] == df[i,"num_likes"] + df[i,"num_loves"] + df[i,"num_wows"] + df[i,"num_hahas"] + df[i,"num_sads"])
  {
    count[i] = TRUE
  }
  else
  {
    count[i] = FALSE
  }
}
```

```

    l[i] = df[i,"num_reactions"] - (df[i,"num_likes"] + df[i,"num_loves"] + df[i,"num_wows"] + df[i,"num_hahas"] + df[i,"num_sads"] + df[i,"num_angrys"])
  }
}
df[count==FALSE,]

```

```

##      num_reactions num_comments num_shares num_likes num_loves num_wows
## 239           885           462           26          659          220           0
## 248           264              2            0          256              2           5
## 249           313              3            0          297              7           6
## 252           247              6            0          234              9           1
## 254           387              3            0          368             16           1
## 255           178              9            0          170              6           0
## 257           270              3            0          256             10           3
## 258           351              4            1          344              6           0
## 294           616           523           21          459          125          21
##      num_hahas num_sads num_angrys
## 239           2           0           0
## 248           0           0           0
## 249           0           0           0
## 252           0           0           0
## 254           0           0           0
## 255           0           0           0
## 257           0           0           0
## 258           0           0           0
## 294           8           0           1

```

```
l[l!=0]
```

```
## [1] 4 1 3 3 2 2 1 1 2
```

```
#We assume that the data that doesn't follow the assumption is incorrect, we check the total percent of
```

```
# As it is less than 1%, we decide to delete it
```

```
(length(l[l!=0])/nrow(df)) * 100
```

```
## [1] 0.1276596
```

```
# Removing the data that does not justify the assumption
```

```
df <- df[!count==FALSE,]
```

```
str(df)
```

```

## 'data.frame': 7041 obs. of 9 variables:
## $ num_reactions: int 529 150 227 111 213 217 503 295 203 170 ...
## $ num_comments : int 512 0 236 0 0 6 614 453 1 9 ...
## $ num_shares : int 262 0 57 0 0 0 72 53 0 1 ...
## $ num_likes : int 432 150 204 111 204 211 418 260 198 167 ...
## $ num_loves : int 92 0 21 0 9 5 70 32 5 3 ...
## $ num_wows : int 3 0 1 0 0 1 10 1 0 0 ...
## $ num_hahas : int 1 0 1 0 0 0 2 1 0 0 ...
## $ num_sads : int 1 0 0 0 0 0 0 0 0 0 ...
## $ num_angrys : int 0 0 0 0 0 0 3 1 0 0 ...

```

```
# Scaling the data for better analysis of clusters
```

```
df <- scale(df)
df<- as.data.frame(df)
head(df)
```

```
##   num_reactions num_comments num_shares   num_likes   num_loves   num_wows
## 1  0.646222253   0.32297476  1.6854261  0.482786321  1.98781086  0.19654943
## 2 -0.172663058  -0.25219846 -0.3042799 -0.144283441 -0.31800076 -0.14742021
## 3 -0.006293219   0.01292045  0.1285950 -0.024206253  0.20832581 -0.03276367
## 4 -0.256928301  -0.25219846 -0.3042799 -0.231005855 -0.31800076 -0.14742021
## 5 -0.036542280  -0.25219846 -0.3042799 -0.024206253 -0.09243223 -0.14742021
## 6 -0.027899691  -0.24545815 -0.3042799 -0.008640691 -0.19268491 -0.03276367
##   num_hahas   num_sads num_angrys
## 1  0.07681282  0.4730465 -0.1556598
## 2 -0.17579767 -0.1526759 -0.1556598
## 3  0.07681282 -0.1526759 -0.1556598
## 4 -0.17579767 -0.1526759 -0.1556598
## 5 -0.17579767 -0.1526759 -0.1556598
## 6 -0.17579767 -0.1526759 -0.1556598
```

```
# Making a vector of status type groups. It can be used in case of verifying cluster accuracy.
```

```
group <- data$status_type

for (i in 1:length(group))
{
  if (group[i] == "video")
  {
    group[i] <- 1
  }
  else if (group[i] == "photo")
  {
    group[i] <- 2
  }
  else if (group[i] == "link")
  {
    group[i] <- 3
  }
  else if (group[i] == "status")
  {
    group[i] <- 4
  }
}
group <- group[!count==FALSE]
summary(group)
```

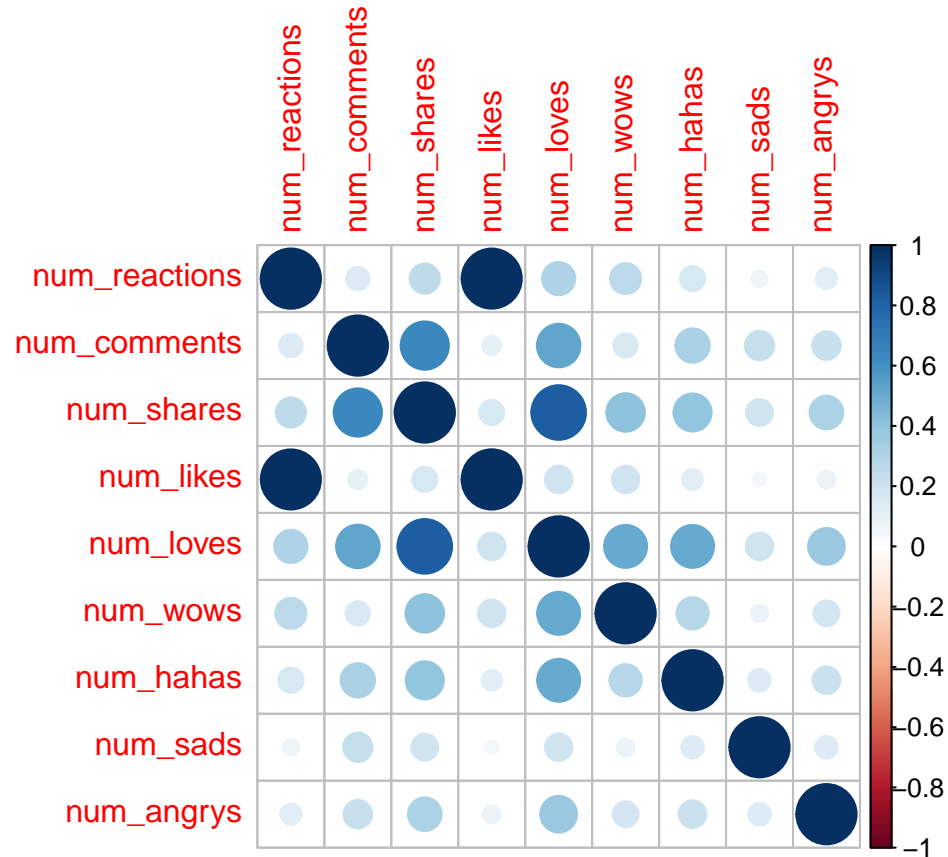
```
##   Length      Class      Mode
##   7041 character character
```

Correlation Matrix

We can that the number of likes and number of reactions have a high correlation; Also, the number of the loves statues and number of shares; the number of loves and number of hahas; number of loves and number

of wows; the number of loves and number of comments; the number of shares and number of comments have a significant correlation.

```
corrplot(cor(df))
```



```
cor(df)
```

```
##          num_reactions num_comments num_shares  num_likes num_loves
## num_reactions    1.0000000    0.1508182  0.2508625  0.99494086  0.3044552
## num_comments     0.1508182    1.0000000  0.6406324  0.10167228  0.5221832
## num_shares       0.2508625    0.6406324  1.0000000  0.17258510  0.8221813
## num_likes        0.99494086    0.1016723  0.1725851  1.00000000  0.2089152
## num_loves        0.30445520    0.5221832  0.8221813  0.20891524  1.0000000
## num_wows         0.26766210    0.1623931  0.4078910  0.20773640  0.5094976
## num_hahas        0.17584823    0.3250046  0.3999405  0.12066387  0.5082269
## num_sads         0.07522229    0.2364420  0.1999311  0.05222869  0.2082756
## num_angrys       0.12427280    0.2251306  0.3125406  0.08739908  0.3715780
##          num_wows num_hahas  num_sads num_angrys
## num_reactions 0.26766210 0.1758482 0.07522229 0.12427280
## num_comments  0.16239310 0.3250046 0.23644201 0.22513056
## num_shares    0.40789100 0.3999405 0.19993111 0.31254065
## num_likes     0.20773640 0.1206639 0.05222869 0.08739908
## num_loves     0.50949759 0.5082269 0.20827560 0.37157800
## num_wows      1.00000000 0.2873832 0.08660209 0.18280525
## num_hahas     0.28738322 1.0000000 0.14148086 0.21165205
```



```
## num_sads      0.08660209 0.1414809 1.00000000 0.14209099
## num_angrys    0.18280525 0.2116520 0.14209099 1.00000000
```

KMeans Clustering

As we have 4 status types - video, photo, status and links, I will start with building a model with 4 clusters.

```
model <- kmeans(df,centers = 4, nstart = 20)
summary(model)
```

```
##           Length Class  Mode
## cluster      7041  -none- numeric
## centers        36  -none- numeric
## totss          1  -none- numeric
## withinss       4  -none- numeric
## tot.withinss   1  -none- numeric
## betweenss      1  -none- numeric
## size           4  -none- numeric
## iter           1  -none- numeric
## ifault         1  -none- numeric
```


Observing the status types in each clusters.

```
table(data[!count == FALSE,]$status_type, model$cluster)
```

```
##
##      1      2      3      4
## link    0    49    14     0
## photo   1 4045   212    23
## status  0   295    70     0
## video  33 1870    76   353
```


Plotting Reactions vs Comments in 4 clusters

```
ggplot(data = df, aes(num_reactions, num_comments, color = model$cluster)) + geom_point()
```



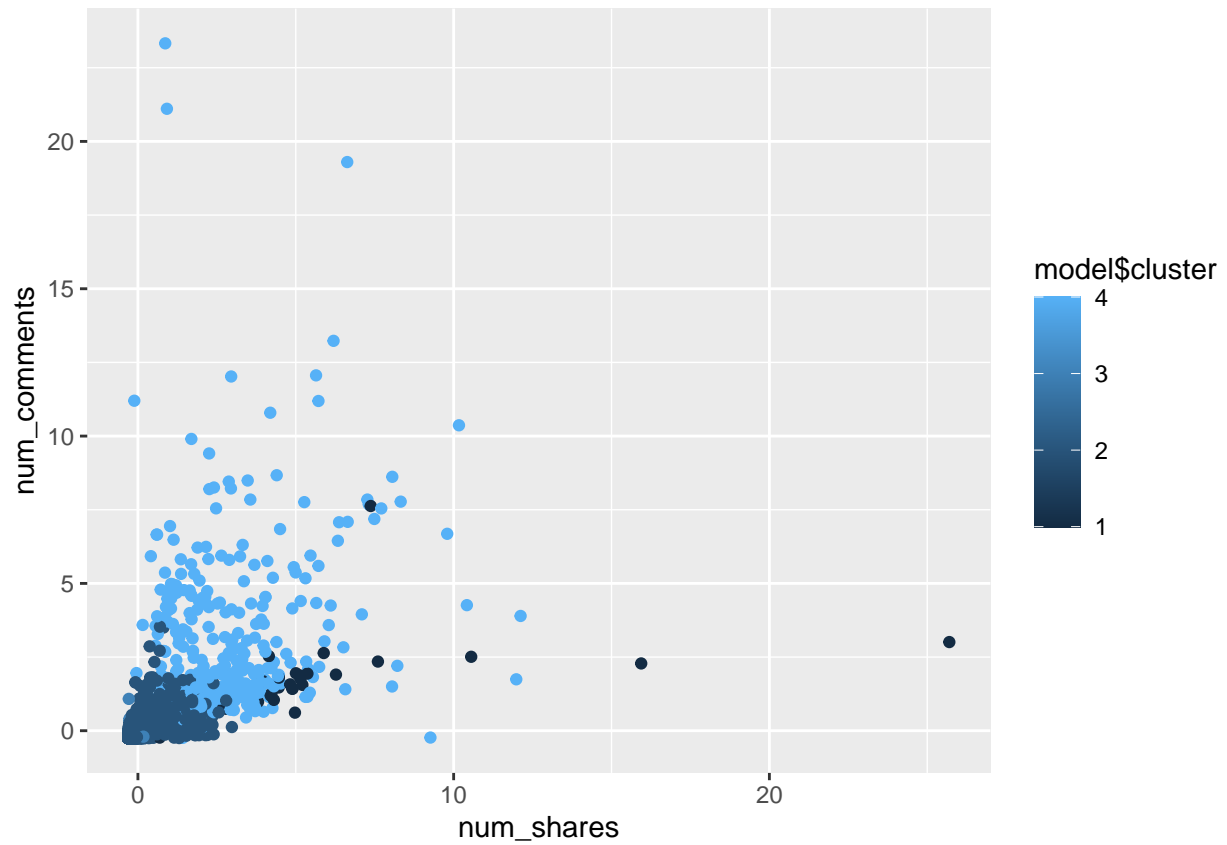
Plotting number of reactions vs number of shares in 4 clusters

```
ggplot(data = df, aes(num_reactions, num_shares, color = model$cluster)) + geom_point()
```




Plotting number of shares vs number of comments in 4 clusters

```
ggplot(data = df, aes(num_shares, num_comments, color = model$cluster)) + geom_point()
```




Plotting positive emotions (love, wow, haha and likes) vs negative emotions (sad and angry) in 4 cluster analysis

```
ggplot(data = df, aes(num_likes+num_loves+num_wows + num_hahas, num_sads+num_angrys, color = model$cluster))
```



Clusplot to plot 2 components in 4 clusters. The first two components explain ~ 58 % of the variability

```
clusplot(df,model$cluster)
```




These two components explain 57.64 % of the point variability.


Finding Optimal Number of Clusters

Using Elbow method, Silhouette method and Gap statistic method I try to calculate the optimal number of clusters. Based on my intuition, I chose 4 earlier as there were 4 groups of status type and it will be good to check if it was the right choice.


```
# Elbow method  
fviz_nbclust(df, kmeans, method = "wss") + labs(subtitle = "Elbow method")
```



```
# Plotting elbow plot with variance explained  
opt <- Optimal_Clusters_KMeans(df, max_clusters=10, plot_clusters = TRUE)
```




```
# Silhouette method  
fviz_nbclust(df, kmeans, method = "silhouette") + labs(subtitle = "Silhouette method")
```

```
# Gap statistic method  
fviz_nbclust(df, kmeans, nstart = 20, method = "gap_stat", nboot = 10) + labs(subtitle = "Gap statistic")
```

```
## Warning: did not converge in 10 iterations  
## Warning: did not converge in 10 iterations  
## Warning: did not converge in 10 iterations  
## Warning: did not converge in 10 iterations  
## Warning: did not converge in 10 iterations  
## Warning: did not converge in 10 iterations  
## Warning: did not converge in 10 iterations
```

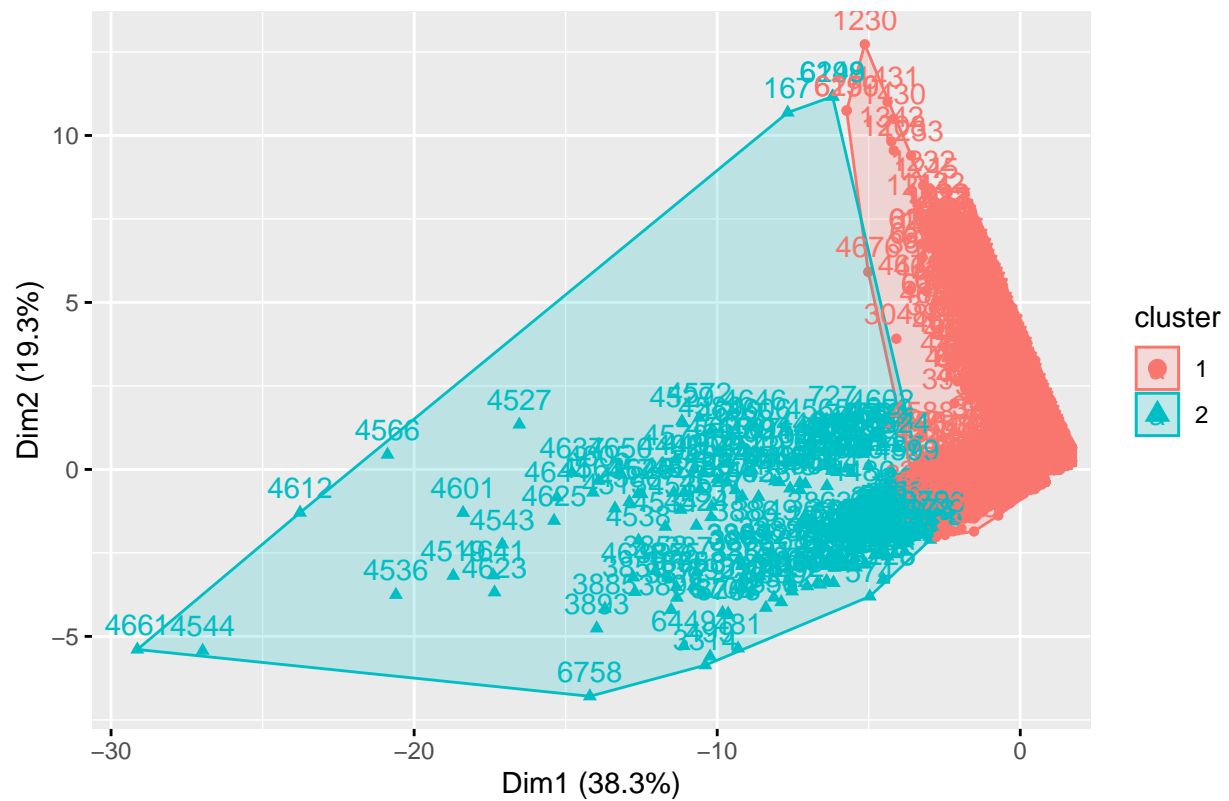


Analysis of $K = 2, 3$ and 4. KMeans & PAM

In K-means, we set number of cluster is 2. The silhouette width is 0.77 and the plot looks convincing.


```
# Investigating k =2  
fviz_cluster(kmeans(df, 2, nstart = 20), data = df)
```

Cluster plot



```
km.sil<-silhouette((kmeans(df,2))$cluster, dist(df))
fviz_silhouette(km.sil)
```


```
## cluster size ave.sil.width
## 1 1 6426 0.86
## 2 2 615 -0.15
```



In K-means, we set number of cluster is 3. The silhouette width is 0.75 which is a bit lower than 2 clusters.


```
# Investigating k =3  
fviz_cluster(kmeans(df, 3, nstart = 20), data = df)
```

Cluster plot



```
km.sil<-silhouette((kmeans(df,3))$cluster, dist(df))
fviz_silhouette(km.sil)
```


```
##  cluster size ave.sil.width
## 1      1  273      -0.11
## 2      2 6377       0.80
## 3      3  391       0.49
```



In K-means, we set number of cluster is 4. The silhouette width is 0.76 which is higher than 3 clusters but lower than 2. It is in the middle and the plot looks convincing.

```
fviz_cluster(kmeans(df, 4, nstart = 20), data = df)
```

Cluster plot




```
km.sil<-silhouette((kmeans(df,4))$cluster, dist(df))
fviz_silhouette(km.sil)
```

##	cluster	size	ave.sil.width
## 1	1	34	0.06
## 2	2	376	-0.05
## 3	3	372	0.51
## 4	4	6259	0.83

Clusters silhouette plot


Average silhouette width: 0.76



Using PAM for additional analysis


```
# Investigating K = 2  
c1<-eclust(df, "pam", k= 2)
```


PAM Clustering




```
fviz_silhouette(c1)
```

```
##   cluster size ave.sil.width
## 1      1 6511         0.80
## 2      2  530         0.04
```




```
fviz_cluster(c1)
```

Cluster plot




```
# Investigating k = 3
c2<-eclust(df, "pam", k= 3)
```

PAM Clustering




```
fviz_silhouette(c2)
```

```
##   cluster size ave.sil.width
## 1      1 803      -0.17
## 2     5824      0.85
## 3      414      0.45
```




```
fviz_cluster(c2)
```

Cluster plot



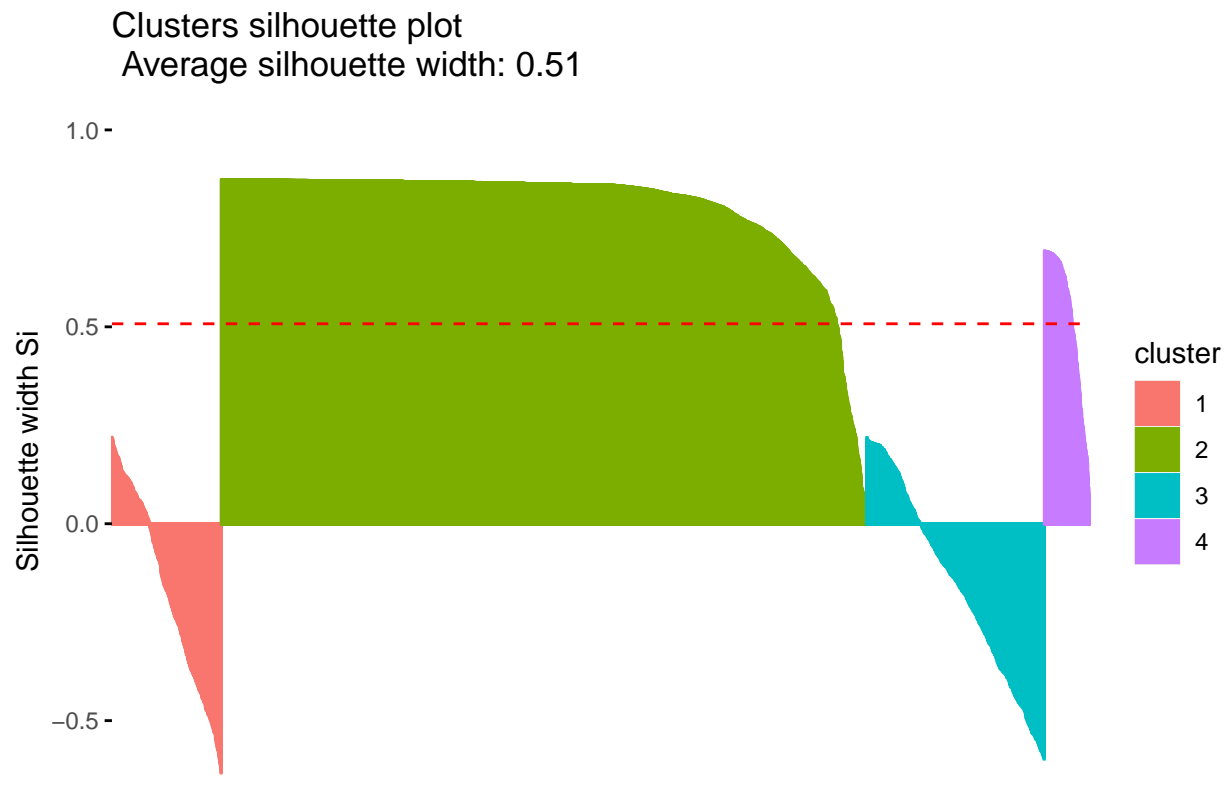
```
# Investigating k = 4
c3<-eclust(df, "pam", k=4)
```

PAM Clustering



```
fviz_silhouette(c3)
```

```
##   cluster size ave.sil.width
## 1      1 786      -0.17
## 2      2 4645      0.80
## 3      3 1277     -0.15
## 4      4  333      0.51
```



```
fviz_cluster(c3)
```


A PCA plot showing the first two principal components, Dim1 (38.3%) and Dim2 (19.3%). The plot displays 1000 samples, each represented by a colored circle. The samples are colored according to their cluster membership, with a legend on the right indicating four clusters: 1 (red), 2 (green), 3 (cyan), and 4 (purple). The plot shows a clear separation between the clusters, with cluster 1 occupying the left side, cluster 2 at the bottom, cluster 3 on the right, and cluster 4 at the top. The axes are labeled with their respective percentages of variance explained.

In the paper we analyzed k-means algorithm on the dataset. We found out that k=2 and k=4 will provide good results. We analysed the results from k-means with PAM clustering. K-means provided better results so we stick to it. As we have 4 status type, (video, photos, statuses, and links) if we try to match the points to cluster prediction and their group vector that I created earlier, for k= 2 it will lead to a lot of errors as there is no group 3 or 4. I will stick with k = 4 and it is confirmed to be a good result with the analysis from k-means.

1. The K-Means Clustering Algorithm Intuition Demonstrated In R, https://uc-r.github.io/kmeans_clustering#elbow
2. University of Warsaw, Unsupervised Learning Course by dr Jacek Lewkowicz